

# 从零开始的RISC-V模拟器开发

## 第11讲 QEMU篇之中断虚拟化1

中国科学院软件研究所  
PLCT实验室

王俊强 wangjunqiang@iscas.ac.cn

李威威 liweiwei@iscas.ac.cn

吴伟 wuwei2016@iscas.ac.cn

# 本课内容

## 中断虚拟化

- QEMU中的IRQ
- ECLIC实现原理

## RISC-V中断

### ➤ 外部中断External Interrupt

- 来自核心外的中断，常见的GPIO、UART中断

### ➤ 定时器中断Timer Interrupt

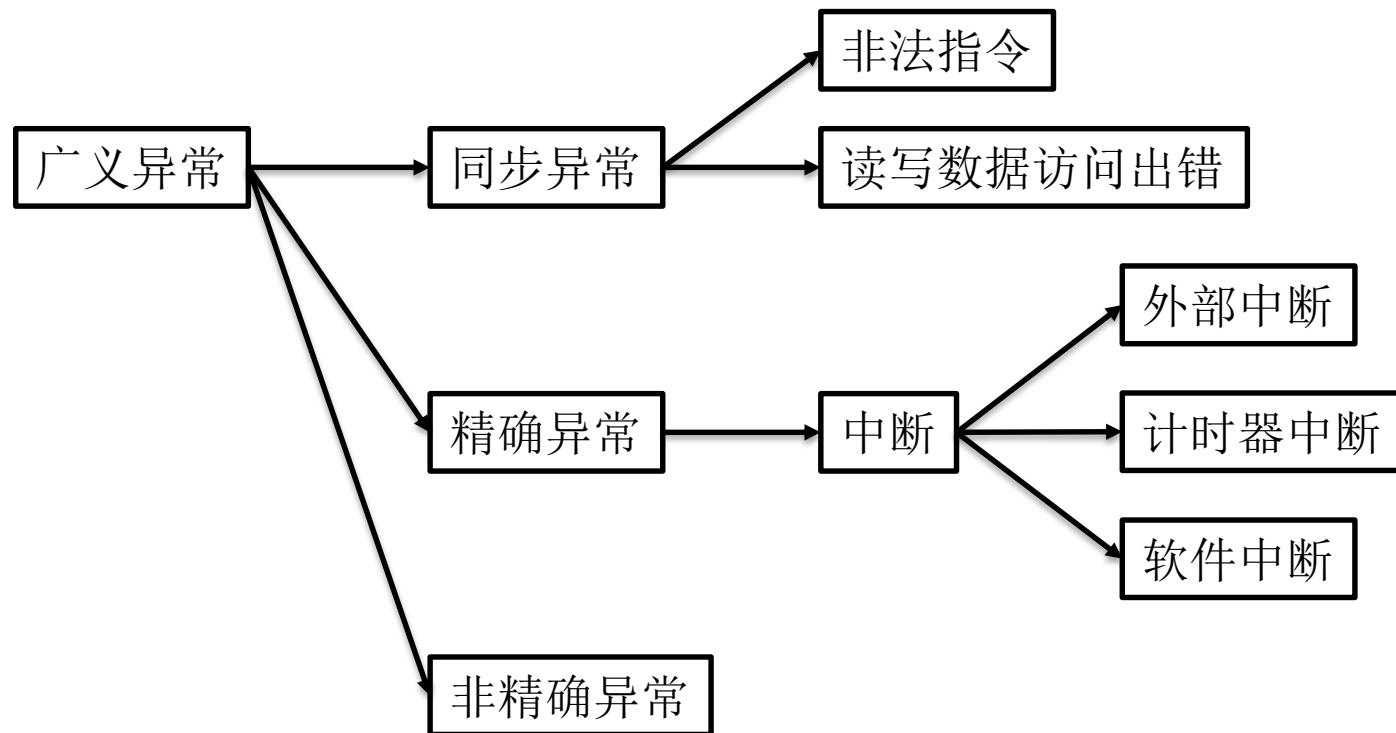
- 来自定时器的中断

### ➤ 软件中断Software Interrupt


- 来自软件自己触发的中断

### ➤ 调试中断Debug Interrupt

- 用于实现调试器的中断



来自:《RISC-V架构与嵌入式开发快速入门》



CPU运行过程中

# QEMU中的IRQ之基础类型

irq类型定义:

```
typedef struct NucLeiUARTState
{
    .....
    /*< public >*/
    qemu_irq irq;
}
```

```
struct SIFIVEGPIOState {
    .....
    qemu_irq irq[SIFIVE_GPIO_PINS];
    qemu_irq output[SIFIVE_GPIO_PIN
S];
    .....
}
```

```
typedef struct NucLeiSYSTIMERState
{
    .....
    DeviceState *eclic;
    qemu_irq *timer_irq;
    qemu_irq *soft_irq;
    .....
}
```

include\qemu\typedefs.h

```
typedef struct IRQState *qemu_irq;
```

```
struct IRQState {
    Object parent_obj;

    qemu_irq_handler handler;
    void *opaque;
    int n;
};
```

n是irq num  
opaque指向所属设备  
qemu\_irq\_handler类型函数指针  
handler回调函数

```
qemu_irq qemu_allocate_irq(qemu_irq_handler handler, void *opaque, int n)
{
    struct IRQState *irq;

    irq = IRQ(object_new(TYPE_IRQ));
    irq->handler = handler;
    irq->opaque = opaque;
    irq->n = n;

    return irq;
}
```

基于QOM的TYPE\_IRQ创建和索引

```
static const TypeInfo irq_type_info = {
    .name = TYPE_IRQ,
    .parent = TYPE_OBJECT,
    .instance_size = sizeof(struct IRQState),
};
```

```
typedef void (*qemu_irq_handler)(void *opaque,
    int n, int level);
```

创建

触发

```
void qemu_set_irq(qemu_irq irq, int level)
{
    if (!irq)
        return;

    irq->handler(irq->opaque, irq->n, level);
}
```

接口:

qemu\_allocate\_irq

qemu\_extend\_irqs

qemu\_allocate\_irqs

qemu\_free\_irqs

qemu\_set\_irq

qemu\_irq\_raise

qemu\_irq\_lower

qemu\_irq\_pulse

# QEMU中的IRQ之基础类型

## Device与GPIO

```
struct DeviceState {  
    .....  
    QLIST_HEAD(, NamedGPIOList) gpios;  
    .....  
};
```

```
struct NamedGPIOList {  
    char *name;  
    qemu_irq *in;    all_in首地址  
    int num_in;  
    int num_out;  
    QLIST_ENTRY(NamedGPIOList) node;  
};
```

hw\core\qdev.c

获取:

qdev\_get\_gpio\_in

qdev\_get\_gpio\_in\_named

```
void qdev_init_gpio_in(DeviceState *dev, qemu_irq_handler handler, int n)
```

qdev\_init\_gpio\_in

qdev\_init\_gpio\_in\_named

qdev\_init\_gpio\_in\_named\_with\_opaque

qemu\_extend\_irqs

object\_property\_add\_child

创建与初始化

```
void qdev_init_gpio_out(DeviceState *dev, qemu_irq *pins, int n)
```

qdev\_init\_gpio\_out

qdev\_init\_gpio\_out\_named

object\_property\_add\_link

连接:

qdev\_connect\_gpio\_out\_named

object\_property\_add\_child

object\_property\_set\_link

# QEMU中的IRQ之处理

target/riscv/cpu.c

```
static struct TCGCPUOps riscv_tcg_ops = {  
    .....  
    .cpu_exec_interrupt = riscv_cpu_exec_interrupt,  
    .....  
    .do_interrupt = riscv_cpu_do_interrupt,  
    .....  
};
```

文件:accel/tcg/cpu-exec.c

```
int cpu_exec(CPUState *cpu)  
{  
    .....  
    cc->tcg_ops->do_interrupt(cpu)  
  
    while (!cpu_handle_exception(cpu, &ret)) {  
        TranslationBlock *last_tb = NULL;  
        int tb_e  
        cc->tcg_ops->cpu_exec_interrupt()  
        while (!cpu_handle_interrupt(cpu, &last_tb)) do_interrupt  
            uint32_t cflags = cpu->cflags_next_tb;  
            TranslationBlock *tb;  
            .....  
        }  
    }  
    cc->cpu_exec_exit(cpu);  
    return ret;  
}
```

exception处理:

```
static inline bool cpu_handle_exception(CPUState *cpu, int *ret)  
{  
    if (cpu->exception_index < 0) {  
        .....  
        return false;  
    }  
    if (cpu->exception_index >= EXCP_INTERRUPT) {  
        .....  
        if (*ret == EXCP_DEBUG) {  
            cpu_handle_debug_exception(cpu);  
        }  
        .....  
    } else {  
        .....  
        if (replay_exception()) {  
            .....  
            qemu_mutex_lock_iothread();  
            cc->tcg_ops->do_interrupt(cpu);  
            qemu_mutex_unlock_iothread();  
            cpu->exception_index = -1;  
            .....  
        }  
        .....  
    }  
    return false;  
}
```

#define EXCP_INTERRUPT	0x10000
#define EXCP_HLT	0x10001
#define EXCP_DEBUG	0x10002
#define EXCP_HALTED	0x10003
#define EXCP_YIELD	0x10004
#define EXCP_ATOMIC	0x10005

## QEMU中的IRQ之处理

target\riscv\cpu.c

```
static struct TCGCPUOps riscv_tcg_ops = {
    .....
    .cpu_exec_interrupt = riscv_cpu_exec_interrupt,
    .....
    .do_interrupt = riscv_cpu_do_interrupt,
    .....
};
```


文件:accel/tcg/cpu-exec.c

```
int cpu_exec(CPUState *cpu)
{
    .....
    cc->tcg_ops->do_interrupt(cpu)
    while (!cpu_handle_exception(cpu, &ret)) {
        TranslationBlock *last_tb = NULL;
        int tb_e
        cc->tcg_ops->cpu_exec_interrupt()
        while (!cpu_handle_interrupt(cpu, &last_tb))
            do_interrupt
            uint32_t cflags = cpu->cflags_next_tb;
            TranslationBlock *tb;
            .....
        }
    }
    cc->cpu_exec_exit(cpu);
    return ret;
}
```

```
static inline bool cpu_handle_interrupt(CPUState *cpu,
                                       TranslationBlock **last_tb)
{
    .....
    if (unlikely(qatomic_read(&cpu->interrupt_request))) {
        interrupt_request = cpu->interrupt_request;
        .....
        if (interrupt_request & CPU_INTERRUPT_DEBUG) {
            cpu->interrupt_request &= ~CPU_INTERRUPT_DEBUG;
            cpu->exception_index = EXCP_DEBUG;
            .....
        }
        if (interrupt_request & CPU_INTERRUPT_HALT) {
            cpu->interrupt_request &= ~CPU_INTERRUPT_HALT;
            cpu->halted = 1;
            cpu->exception_index = EXCP_HLT;
            .....
            return true;
        }
        else if (interrupt_request & CPU_INTERRUPT_RESET) {
            cpu_reset(cpu);
            .....
            return true;
        }
        else {
            if (cc->tcg_ops->cpu_exec_interrupt &&
                cc->tcg_ops->cpu_exec_interrupt(cpu, interrupt_request)) {
                .....
            }
        }
        if (interrupt_request & CPU_INTERRUPT_EXITTB) {
            .....
        }
    }
    cpu->interrupt_request!
}
```

interrupt处理

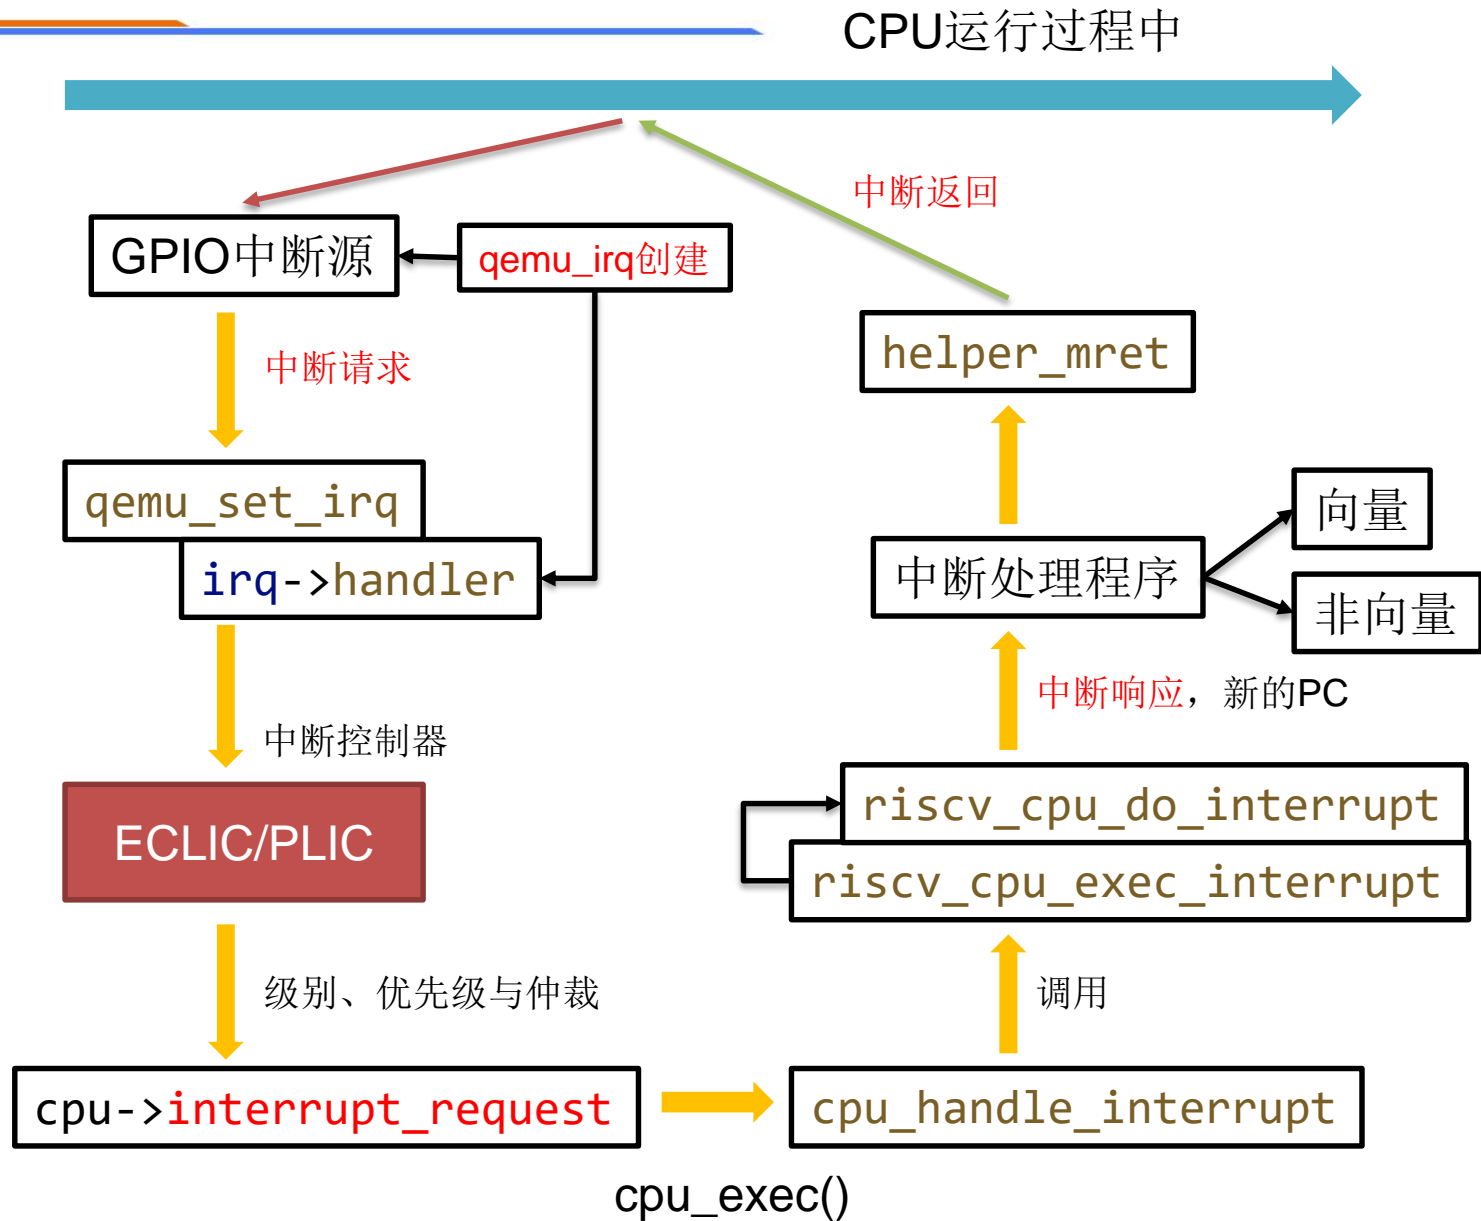
## NUCLEI IRQ主要流程



CPU运行过程中


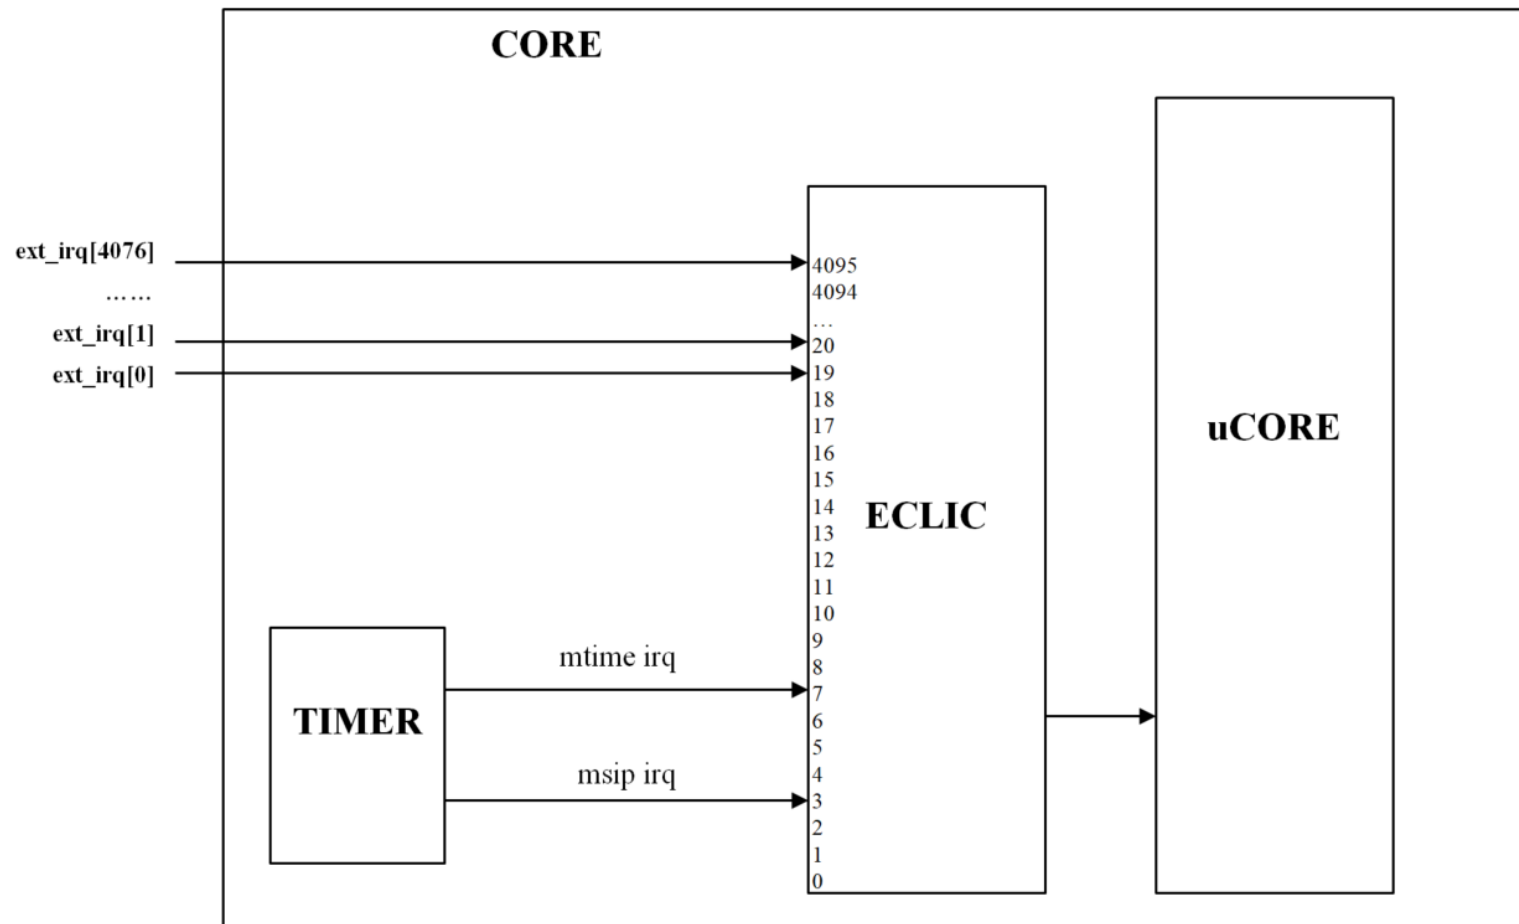
主要目标:

- qemu\_irq创建
- ECLIC实现
- riscv\_cpu\_exec\_interrupt更新
- riscv\_cpu\_do\_interrupt更新
- 相应CSR的实现
- helper\_mret更新





# NUCLEI ECLIC(Enhanced Core Local Interrupt Controller)

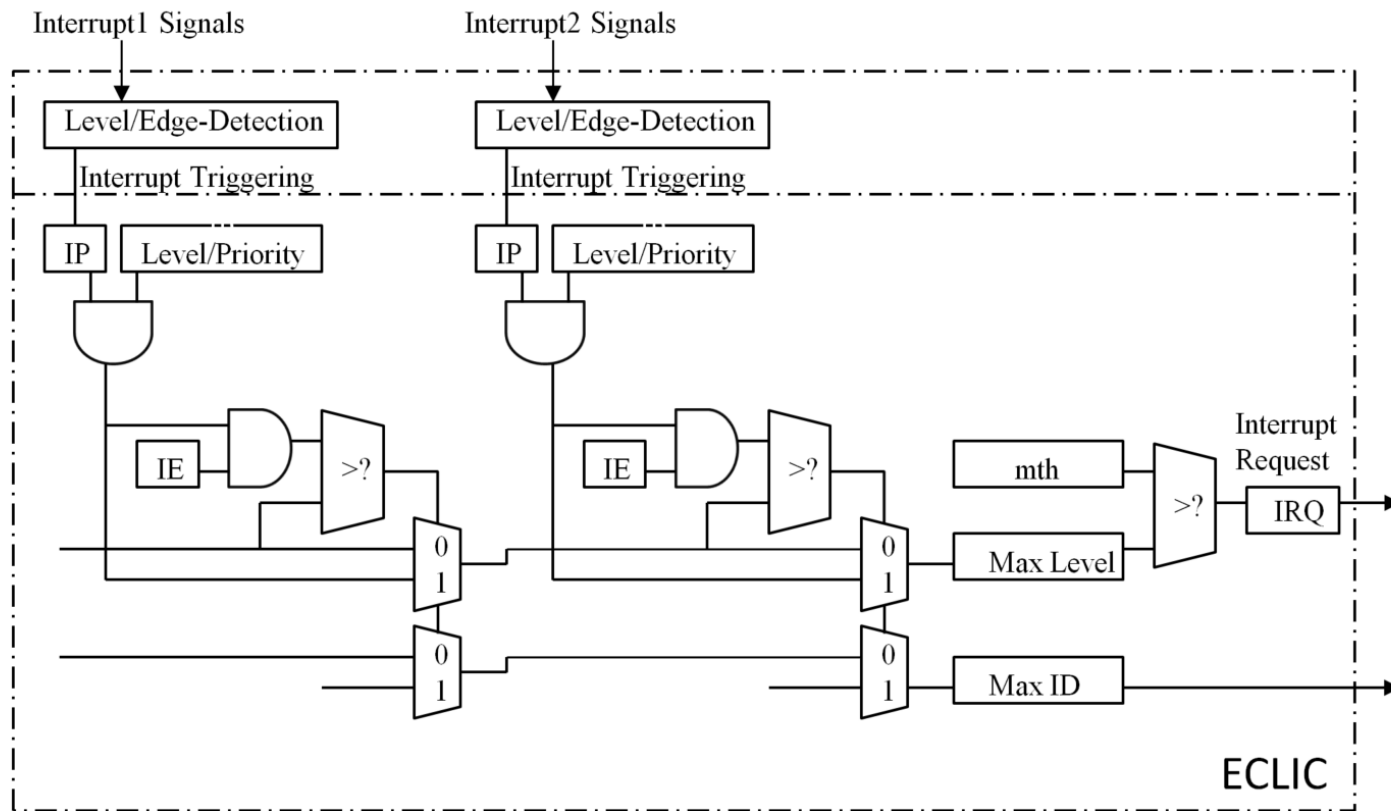


Machine mode

SYSTIMER+ ECLIC/CLIC for mcu  
CLINT+PLIC for linux

- ECLIC只服务于一个处理器内核，为该处理器内核私有
- ECLIC的软件编程模型也向后兼容标准的 CLIC

# NUCLEI ECLIC(Enhanced Core Local Interrupt Controller)



## 相关概念:

- 编号ID
- 使能位IE
- 等待标志位IP
- 电平或边沿属性 Level or Edge Triggered
- 级别和优先级 Level and Priority
- 向量或非向量处理 Vector or Non Vector Mode

阅读: Nuclei\_N级别指令架构手册.pdf

## QEMU中的IRQ之基础类型

### ECLIC Registers:

Offset	Permission	Register	Width
0x0000	RW	cliccfg	8-bit
0x0004	R	clicinfo	32-bit
0x000b	RW	mth	8-bit
0x1000+4*i	RW	clicintip[i]	8-bit
0x1001+4*i	RW	clicintie[i]	8-bit
0x1002+4*i	RW	clicintattr[i]	8-bit
0x1003+4*i	RW	clicintctl[i]	8-bit

来自:Nuclei\_N级别指令架构手册.pdf

include/hw/intc/nuclei\_eclic.h

```
typedef struct NucLeiECLICState {
    /*< private >*/
    SysBusDevice parent_obj;
    /*< public >*/
    MemoryRegion mmio;
    uint32_t num_sources;
    /* config */
    uint32_t sources_id;
    uint8_t cliccfg;
    uint32_t clicinfo;
    uint8_t mth;
    uint8_t *clicintip;
    uint8_t *clicintie;
    uint8_t *clicintattr;
    uint8_t *clicintctl;
    ECLICPendingInterrupt *clicintlist;
    uint32_t aperture_size;

    QLIST_HEAD(, ECLICPendingInterrupt)
    pending_list;
    size_t active_count;

    /* ECLIC IRQ handlers */
    qemu_irq *irqs;
} NucLeiECLICState;
```

# NUCLEI ECLIC之创建

eclic的静态注册:

```
static const TypeInfo nuclei_eclic_info = {  
    .name = TYPE_NUCLEI_ECLIC,  
    .parent = TYPE_SYS_BUS_DEVICE,  
    .instance_size = sizeof(NucLeiECLICState),  
    .class_init = nuclei_eclic_class_init,  
};  
  
static void nuclei_eclic_register_types(void)  
{  
    type_register_static(&nuclei_eclic_info);  
}  
  
type_init(nuclei_eclic_register_types);
```

mmio接口设置:

```
static const MemoryRegionOps nuclei_eclic_ops =  
{  
    .read = nuclei_eclic_read,  
    .write = nuclei_eclic_write,  
    .endianness = DEVICE_LITTLE_ENDIAN,  
};
```

```
static void nuclei_eclic_realize(DeviceState *dev, Error **errp)  
{  
    NucLeiECLICState *eclic = NUCLEI_ECLIC(dev);  
    int id;  
  
    memory_region_init_io(&eclic->mmio, OBJECT(dev), &nuclei_eclic_ops, eclic,  
                          TYPE_NUCLEI_ECLIC, eclic->aperture_size);  
    sysbus_init_mmio(SYS_BUS_DEVICE(dev), &eclic->mmio);  
  
    eclic->clicintip = g_new0(uint8_t, eclic->num_sources);  
    .....  
    QLIST_INIT(&eclic->pending_list);  
    for (id = 0; id < eclic->num_sources; id++) {  
        eclic->clicintlist[id].irq = id;  
        update_eclic_int_info(eclic, id);  
    }  
    eclic->active_count = 0;  
  
    /* Init ECLIC IRQ */  
    eclic->irqs[Internal_SysTimerSW_IRQn] =  
        qemu_allocate_irq(nuclei_eclic_irq_request,  
                          eclic, Internal_SysTimerSW_IRQn);  
    eclic->irqs[Internal_SysTimer_IRQn] =  
        qemu_allocate_irq(nuclei_eclic_irq_request,  
                          eclic, Internal_SysTimer_IRQn);  
  
    for (id = Internal_Reserved_Max_IRQn; id < eclic->num_sources; id++) {  
        eclic->irqs[id] = qemu_allocate_irq(nuclei_eclic_irq_request,  
                                             eclic, id);  
    }  
    .....  
}
```

# NUCLEI ECLIC之读写

## read and write

```
static uint64_t nuclei_ecllic_read(void *opaque, hwaddr offset,
unsigned size)
{
    NucLeiECLICState *eclic = NUCLEI_ECLIC(opaque);
    .....
    switch (offset) {
    case NUCLEI_ECLIC_REG_CLICCFG:
        value = eclic->cllicfg & 0xFF;
        break;
    case NUCLEI_ECLIC_REG_CLICINFO:
        value = (CLICINTCTLBITS << 21) & 0xFFFFFFFF;
        break;
    case NUCLEI_ECLIC_REG_MTH:
        value = eclic->mth & 0xFF;
        break;
    case NUCLEI_ECLIC_REG_CLICINTIP_BASE:
        value = eclic->clicintip[id] & 0xFF;
        break;
    case NUCLEI_ECLIC_REG_CLICINTIE_BASE:
        value = eclic->clicintie[id] & 0xFF;
        break;
    .....
    }
    return value;
}
```

```
static void nuclei_ecllic_write(void *opaque, hwaddr offset, uint64_t value,
unsigned size)
{
    NucLeiECLICState *eclic = NUCLEI_ECLIC(opaque);
    uint32_t id = 0;
    .....
    switch (offset) {
    case NUCLEI_ECLIC_REG_CLICCFG:
        .....
    case NUCLEI_ECLIC_REG_MTH:
        nuclei_ecllic_update_intmth(eclic, id, value & 0xFF);
        break;
    case NUCLEI_ECLIC_REG_CLICINTIP_BASE:
        if ((eclic->clicintlist[id].trigger & 0x1) != 0) {
            if ((eclic->clicintip[id] == 0) && (value & 0x1) == 1) {
                eclic->clicintip[id] = 1;
                eclic_insert_pending_list(eclic, id);
            } else if ((eclic->clicintip[id] == 1) && (value & 0x1) == 0) {
                eclic->clicintip[id] = 0;
                eclic_remove_pending_list(eclic, id);
            }
        }
        nuclei_ecllic_next_interrupt(eclic);
        break;
    case NUCLEI_ECLIC_REG_CLICINTIE_BASE:
        nuclei_ecllic_update_intie(eclic, id, value & 0xFF);
        break;
    .....
    default:
        break;
    }
}
```

## NUCLEI ECLIC之核心设置函数

```
static void nuclei_eclic_next_interrupt(void *eclic_ptr)
{
    RISCVCPU *cpu = RISCV_CPU(qemu_get_cpu(0));
    CPURISCVState *env = &cpu->env;
    NucLeiECLICState *eclic = (NucLeiECLICState *)eclic_ptr;
    ECLICPendingInterrupt *active;
    target_ulong mil;
    int shv;

    QLIST_FOREACH(active, &eclic->pending_list, next)
    {
        if (active->enable) {
            mil = get_field(env->mintstatus, MINTSTATUS_MIL);
            if (active->level >= eclic->mth && active->level > mil) {
                shv = eclic->clicintattr[active->irq] & 0x1;
                eclic->active_count++;
                riscv_cpu_eclic_interrupt(cpu,
                    (active->irq & 0xFFF) | (shv << 12) | (active->level << 13));
                return;
            }
        }
    }
}
```

cpu->interrupt\_request = CPU\_INTERRUPT\_ECLIC

```
static void riscv_cpu_eclic_interrupt(RISCVCPU *cpu,
                                     int exccode)
{
    CPURISCVState *env = &cpu->env;
    bool locked = false;

    env->exccode = exccode;

    if (!qemu_mutex_iothread_locked()) {
        locked = true;
        qemu_mutex_lock_iothread();
    }

    if (exccode != -1) {
        env->irq_pending = true;
        cpu_interrupt(CPU(cpu), CPU_INTERRUPT_ECLIC);
    } else {
        env->irq_pending = false;
        cpu_reset_interrupt(CPU(cpu), CPU_INTERRUPT_ECLIC);
    }

    if (locked) {
        qemu_mutex_unlock_iothread();
    }
}
```

ECLIC => CPU Request

# NUCLEI ECLIC之riscv\_cpu\_exec\_interrupt

```
bool riscv_cpu_exec_interrupt(CPUState *cs, int interrupt_request)
{
    #if !defined(CONFIG_USER_ONLY)
        .....

        if (interrupt_request & CPU_INTERRUPT_ECLIC) {
            RISCVCPU *cpu = RISCV_CPU(cs);
            CPURISCVState *env = &cpu->env;
            int mode = PRV_M;
            int enabled = riscv_cpu_local_irq_mode_enabled(env, mode);
            if (enabled) {
                cs->exception_index = RISCV_EXCP_INT_ECLIC | env->exccode;
                cs->interrupt_request &= ~CPU_INTERRUPT_ECLIC;
                riscv_cpu_do_interrupt(cs);
                return true;
            }
        }
    #endif
    return false;
}
```

阅读:手册7.4.7

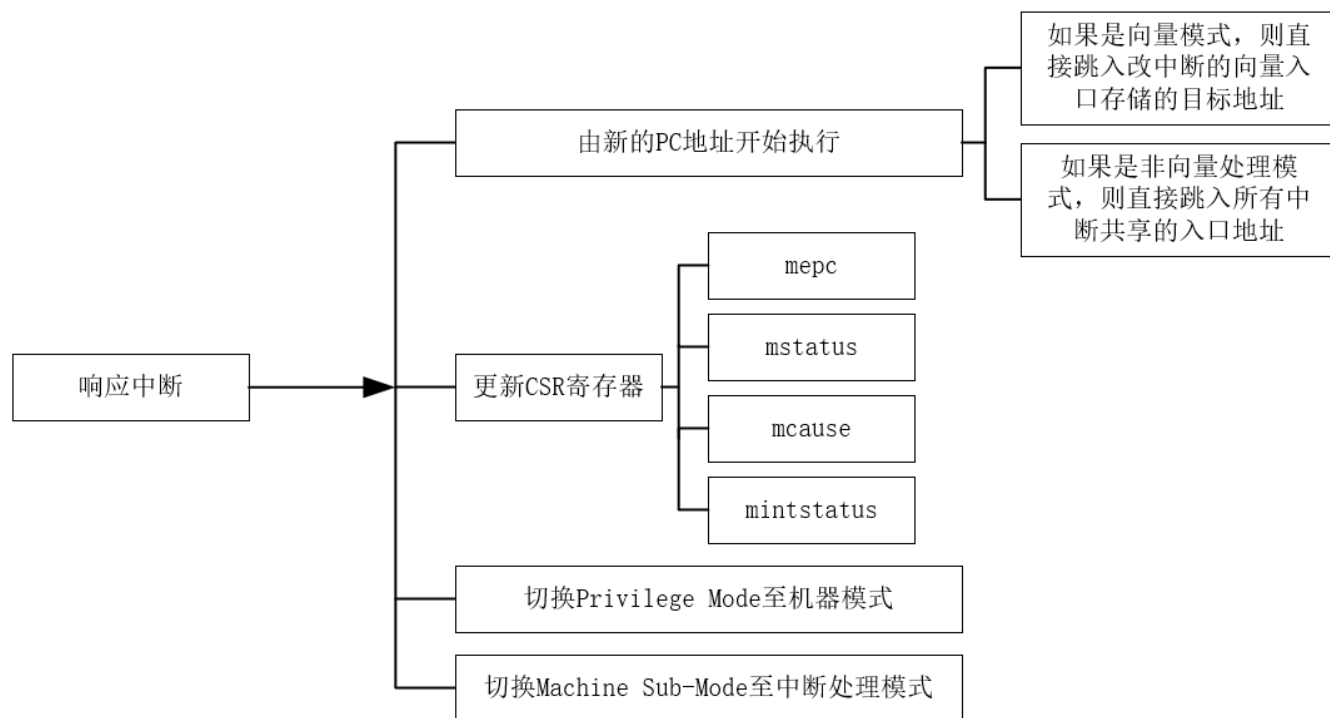
```
#define CPU_INTERRUPT_HARD      0x0002
#define CPU_INTERRUPT_EXITTB    0x0004
/* Halt the CPU. */
#define CPU_INTERRUPT_HALT      0x0020
/* Debug event pending. */
#define CPU_INTERRUPT_DEBUG     0x0080
/* Reset signal. */
#define CPU_INTERRUPT_RESET     0x0400
/* Several target-
specific external hardware interrupts. Each target/cpu.h
should define proper names based on these defines. */
#define CPU_INTERRUPT_TGT_EXT_0 0x0008
#define CPU_INTERRUPT_TGT_EXT_1 0x0010
#define CPU_INTERRUPT_TGT_EXT_2 0x0040
#define CPU_INTERRUPT_TGT_EXT_3 0x0200
#define CPU_INTERRUPT_TGT_EXT_4 0x1000
```

```
static int riscv_cpu_local_irq_mode_enabled(CPURISCVState *env,
int mode)
{
    switch (mode) {
        case PRV_M:
            return env->priv < PRV_M ||
                (env->priv == PRV_M &&
                 get_field(env->mstatus, MSTATUS_MIE));

        .....
        default:
            return false;
    }
}
```

MSTATUS—MIE

# NUCLEI ECLIC之riscv\_cpu\_do\_interrupt



阅读手册： 5.6

mepc:7.4.20

mstatus:7.4.7

mcause: 7.4.21

mintstatus: 7.4.25

向量模式

5.13

非向量模式



# NUCLEI ECLIC之riscv\_cpu\_do\_interrupt

更新寄存器:


```
if(eclic_flag)
{
    mode = (cause >> 12) & 0x1;
    level = (cause >> 13) & 0xFF;
    cause &= 0x3ff;

    cause |= get_field(env->mstatus, MSTATUS_MPP) << 28;
    cause |= get_field(env->mintstatus, MINTSTATUS_MIL) << 16;
    cause |= get_field(env->mstatus, MSTATUS_MPIE) << 27;
    cause = set_field(cause, MCAUSE_MPP, PRV_M);
    cause = set_field(cause, MCAUSE_INTERRUPT, 1);

    env->mintstatus = set_field(env->mintstatus, MINTSTATUS_MIL, level);
}
```

更新PC:

```
env->pc = riscv_intr_pc(env, async, eclic_flag, cause, mode);
```



注意更新MTVEC实现

阅读手册: 7.4.17  
7.5.14

```
if(eclic_flag ) {
    if(mode)
    {
        向量模式下
        uint64_t vec_addr = (cause & 0x3FF) *4 + env->mtvt;
        cpu_physical_memory_rw(vec_addr, &newpc, 4, 0);
    }else{
        if ((env->mtvt2 & 0x1) == 0) {
            newpc = env->mtvec & 0xffffffffc;
        } else if ((env->mtvt2 & 0x1) == 1) {
            newpc = env->mtvt2 & 0xffffffffc;
        }
    }
    非向量模式下
}
```

# NUCLEI ECLIC之中断处理程序

SoC\demosoc\Common\Source\GCC\startup\_demosoc.S

```
/*
 * Intialize ECLIC vector interrupt
 * base address mtvt to vector_base
 */
la t0, vector_base    中断向量表
csrw CSR_MTVT, t0
/*
 * Set ECLIC non-vector entry to be controlled
 * by mtvt2 CSR register.
 * Intialize ECLIC non-vector interrupt
 * base address mtvt2 to irq_entry.
 */
la t0, irq_entry      中断共享入口表
csrw CSR_MTVT2, t0
csrs CSR_MTVT2, 0x1
/*
 * Set Exception Entry MTVEC to exc_entry
 * Due to settings above, Exception and NMI
 * will share common entry.
 */
la t0, exc_entry
csrw CSR_MTVEC, t0
/* Set the interrupt processing mode to ECLIC mode */
li t0, 0x3f
csrc CSR_MTVEC, t0
csrs CSR_MTVEC, 0x3
```

```
vector_base:
#if defined(DOWNLOAD_MODE) && (DOWNLOAD_MODE != DOWNLOAD_MODE_FLASH)
    j _start
    .align LOG_REGBYTES
#else
    DECLARE_INT_HANDLER    default_intexc_handler
#endif
    DECLARE_INT_HANDLER    default_intexc_handler
    DECLARE_INT_HANDLER    default_intexc_handler
    DECLARE_INT_HANDLER    eclic_msip_handler

    DECLARE_INT_HANDLER    default_intexc_handler
    DECLARE_INT_HANDLER    default_intexc_handler
    DECLARE_INT_HANDLER    default_intexc_handler
    DECLARE_INT_HANDLER    eclic_mtip_handler
.....
```

向量模式

```
.global irq_entry
.weak irq_entry
irq_entry:
    SAVE_CONTEXT
    SAVE_CSR_CONTEXT
    /* This special CSR read/write operation, which is actually
     * claim the CLIC to find its pending highest ID, if the ID
     * is not 0, then automatically enable the mstatus.MIE, and
     * jump to its vector-entry-label, and update the link register
     */
    csrrw ra, CSR_JALMNXTI, ra
    DISABLE_MIE
    RESTORE_CSR_CONTEXT
    RESTORE_CONTEXT
    mret
```

非向量模式

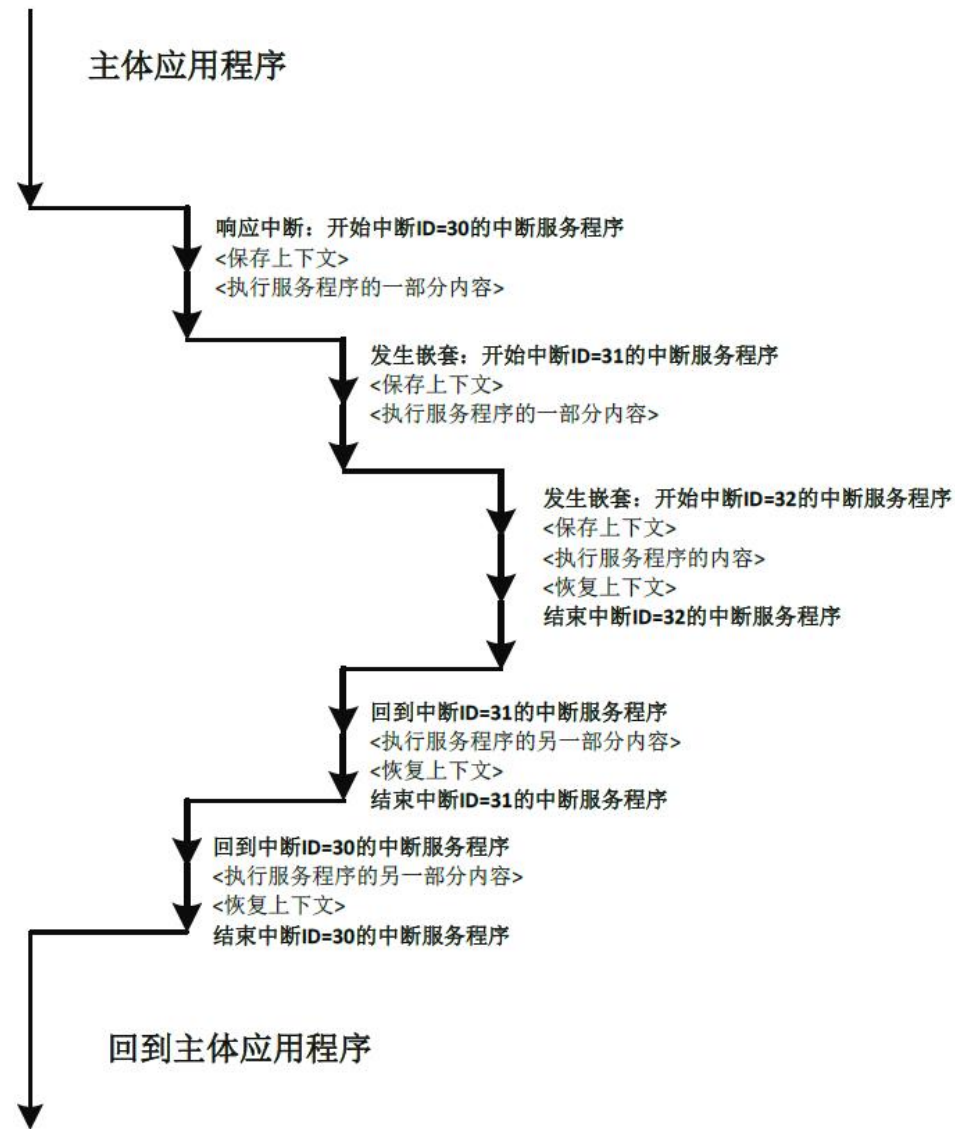
# NUCLEI ECLIC之中断处理程序

application\baremetal\demo\_eclic\demo\_eclic.c

```
__INTERCEPT void eclic_msip_handler(void)
{
    static uint32_t int_sw_cnt = 0;    /* software interrupt counter */
    /
    // save CSR context
    SAVE_IRQ_CSR_CONTEXT();
    SysTimer_ClearSWIRQ();
    printf("[IN SOFTWARE INTERRUPT]software interrupt hit %d times\r\n", int_sw_cnt++);
    printf("[IN SOFTWARE INTERRUPT]software interrupt end\r\n");
    // restore CSR context
    RESTORE_IRQ_CSR_CONTEXT();
}
```

```
#define SAVE_IRQ_CSR_CONTEXT()
    rv_csr_t __mcause = __RV_CSR_READ(CSR_MCAUSE);
    rv_csr_t __mepc = __RV_CSR_READ(CSR_MEPC);
    rv_csr_t __msubm = __RV_CSR_READ(CSR_MSUBM);
    __enable_irq();
```

向量模式下的嵌套  
阅读手册:5.13.2



# NUCLEI ECLIC之中断处理程序

SoC\demosoc\Common\Source\GCC\intexc\_demosoc.S

```
/**
 * \brief Macro for save necessary CSRs to stack
 * \details
 * This macro store MCAUSE, MEPC, MSUBM to stack.
 */
.macro SAVE_CSR_CONTEXT
    /* Store CSR mcause to stack using pushmcause */
    csrrwi x0, CSR_PUSHMCAUSE, 11
    /* Store CSR mepc to stack using pushmepc */
    csrrwi x0, CSR_PUSHMEPC, 12
    /* Store CSR msub to stack using pushmsub */
    csrrwi x0, CSR_PUSHMSUBM, 13
.endm
```

CSR\_JALMNXTI

阅读手册：7.5.13~

非向量的嵌套和咬尾：5.13

```
// If define SWIRQ_INTLEVEL_HIGHER equals 1 the software interrupt
// will have a higher interrupt level.
// the software interrupt will run during timer interrupt.
// If define SWIRQ_INTLEVEL_HIGHER equals 0 the software interrupt
// will have a lower interrupt level.
// the software interrupt will run after timer interrupt.
#define SWIRQ_INTLEVEL_HIGHER 0

// timer interrupt handler
// non-vector mode interrupt
void eclic_mtip_handler(void)
{
    static uint32_t int_t_cnt = 0;
    printf("-----\r\n");
    printf("[IN TIMER INTERRUPT]timer interrupt hit %d times\r\n", in
t_t_cnt++);

    printf("[IN TIMER INTERRUPT]trigger software interrupt\r\n");
    #if SWIRQ_INTLEVEL_HIGHER == 1
        printf("[IN TIMER INTERRUPT]software interrupt will run during ti
mer interrupt\r\n");
    #else
        printf("[IN TIMER INTERRUPT]software interrupt will run when time
r interrupt finished\r\n");
    #endif
    // trigger software interrupt
    SysTimer_SetSWIRQ();
    // Reload Timer Interrupt
    SysTick_Reload(TIMER_TICKS);
    printf("[IN TIMER INTERRUPT]timer interrupt end\r\n");
}
```

# NUCLEI ECLIC之中断处理程序

SoC\demosoc\Common\Source\GCC\intexc\_demosoc.S

```
/**
 * \brief Macro for save necessary CSRs to stack
 * \details
 * This macro store MCAUSE, MEPC, MSUBM to stack.
 */
.macro SAVE_CSR_CONTEXT
    /* Store CSR mcause to stack using pushmcause */
    csrrwi x0, CSR_PUSHMCAUSE, 11
    /* Store CSR mepc to stack using pushmepc */
    csrrwi x0, CSR_PUSHMEPC, 12
    /* Store CSR msub to stack using pushmsub */
    csrrwi x0, CSR_PUSHMSUBM, 13
.endm
```

## CSR\_JALMNXTI

```
static int rmw_pushmcause(CPURISCVState *env, int csrno, target_ulong *ret_value,
                          target_ulong new_value, target_ulong write_mask)
{
    uint64_t notify_addr = new_value * 4 + env->gpr[2];
    cpu_physical_memory_rw(notify_addr, &env->mcause, 4, 1);
    return 0;
}

static int rmw_pushmepc(CPURISCVState *env, int csrno, target_ulong *ret_value,
                        target_ulong new_value, target_ulong write_mask)
{
    uint64_t notify_addr = new_value * 4 + env->gpr[2];
    cpu_physical_memory_rw(notify_addr, &env->mepc, 4, 1);
    return 0;
}

static int rmw_jalmnxti(CPURISCVState *env, int csrno, target_ulong *ret_value,
                       target_ulong new_value, target_ulong write_mask)
{
    target_ulong addr;

    if (env->irq_pending) {
        uint64_t vec_addr = (env->mcause & 0x3FF) * 4 + env->mtvt;
        cpu_physical_memory_rw(vec_addr, &addr, 4, 0);
        env->gpr[1] = env->pc + 4;
        env->gpr[5] = env->pc + 4;
        *ret_value = addr;
        riscv_cpu_eclic_clean_pending(env->eclic, env->mcause & 0x3ff);
        env->mstatus = set_field(env->mstatus, MSTATUS_MIE, 1);
    } else
        *ret_value = env->pc + 4;
    return 0;
}
```

# NUCLEI ECLIC之中断处理程序

SoC\demosoc\Common\Source\GCC\intexc\_demosoc.S

```
/**
 * \brief Macro for save necessary CSRs to stack
 * \details
 * This macro store MCAUSE, MEPC, MSUBM to stack.
 */
.macro SAVE_CSR_CONTEXT
    /* Store CSR mcause to stack using pushmcause */
    csrrwi x0, CSR_PUSHMCAUSE, 11
    /* Store CSR mepc to stack using pushmepc */
    csrrwi x0, CSR_PUSHMEPC, 12
    /* Store CSR msub to stack using pushmsub */
    csrrwi x0, CSR_PUSHMSUBM, 13
.endm
```

```
static bool trans_csrrw(DisasContext *ctx, arg_csrrw *a)
{
    TCGv source1, csr_store, dest, rs1_pass;
    RISCV_OP_CSR_PRE;
    gen_helper_csrrw(dest, cpu_env, source1, csr_store);
    if(a->csr == CSR_JALMNXTI)
        RISCV_OP_CSR_JAL_POST;
    else
        RISCV_OP_CSR_POST;
    return true;
}
```

target\riscv\csr.c

```
static int rmw_jalmnxti(CPURISCVState *env, int csrno, target_ulong *ret_value,
                        target_ulong new_value, target_ulong write_mask)
{
    target_ulong addr;

    if (env->irq_pending) {
        uint64_t vec_addr = (env->mcause & 0x3FF) * 4 + env->mtvt;
        cpu_physical_memory_rw(vec_addr, &addr, 4, 0);
        env->gpr[1] = env->pc + 4;
        env->gpr[5] = env->pc + 4;
        *ret_value = addr;
        riscv_cpu_eclic_clean_pending(env->eclic, env->mcause & 0x3ff);
        env->mstatus = set_field(env->mstatus, MSTATUS_MIE, 1);
    } else
        *ret_value = env->pc + 4;
    return 0;
}
```

```
#define RISCV_OP_CSR_POST do {\
    gen_set_gpr(a->rd, dest); \
    tcg_gen_movi_tl(cpu_pc, ctx->pc_succ_insn); \
    exit_tb(ctx); \
    ctx->base.is_jmp = DISAS_NORETURN; \
    tcg_temp_free(source1); \
    tcg_temp_free(csr_store); \
    tcg_temp_free(dest); \
    tcg_temp_free(rs1_pass); \
} while (0)
```

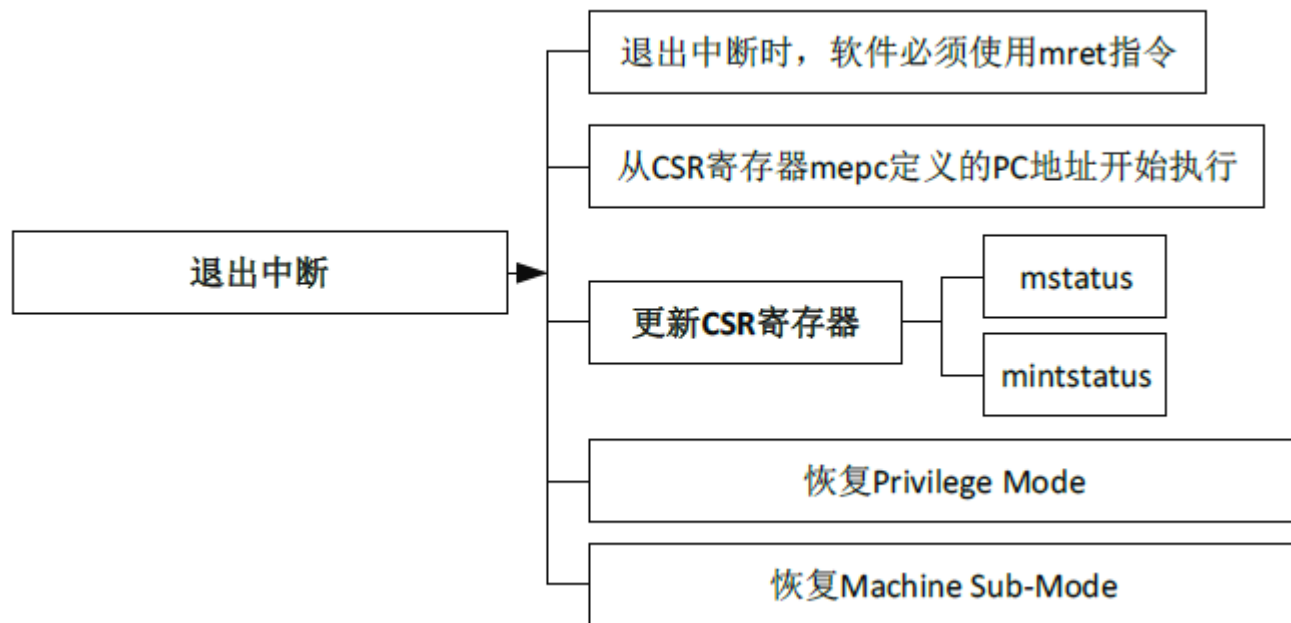
target\riscv\insn\_trans\trans\_rvi.c.inc

咬尾自跳转

CSR\_JALMNXTI

```
#define RISCV_OP_CSR_JAL_POST do {\
    tcg_gen_mov_tl(cpu_pc, dest); \
    exit_tb(ctx); \
    ctx->base.is_jmp = DISAS_NORETURN; \
    tcg_temp_free(source1); \
    tcg_temp_free(csr_store); \
    tcg_temp_free(dest); \
    tcg_temp_free(rs1_pass); \
} while (0)
```

## NUCLEI ECLIC之mret



参考手册：5.7

```
target_ulong helper_mret(CPURISCVState *env,
                        target_ulong cpu_pc_deb)
{
    .....
    target_ulong retpc = env->mepc;
    .....
    target_ulong prev_priv = get_field(mstatus, MSTATUS_MPP);
    mstatus = set_field(mstatus, MSTATUS_MIE,
                        get_field(mstatus, MSTATUS_MPIE));
    mstatus = set_field(mstatus, MSTATUS_MPIE, 1);
    mstatus = set_field(mstatus, MSTATUS_MPP, PRV_U);
    mstatus = set_field(mstatus, MSTATUS_MPV, 0);
    env->mstatus = mstatus;
    riscv_cpu_set_mode(env, prev_priv);
    .....
    if ((env->mtvec & 0b11111) == 0b000011) {
        target_ulong mpil = get_field(env->mcause, MCAUSE_MPIL);
        env->mintstatus = set_field(env->
>mintstatus, MINTSTATUS_MIL, mpil);

        qemu_mutex_lock_iothread();
        riscv_cpu_eclic_get_next_interrupt(env->eclic);
        qemu_mutex_unlock_iothread();

        if (get_field(env->mcause, MCAUSE_INTERRUPT) == 1) {
            env->mstatus = set_field(env->mstatus, MSTATUS_MPP,
                                    get_field(env->mcause, MCAUSE_MPP));
        }
    }
    return retpc;
}
```



# NUCLEI ECLIC之使用

## nuclei\_n\_soc\_realize函数中

hw\riscv\nuclei\_n.c

```
s->eclic = nuclei_eclic_create(memmap[NUCLEI_N_ECLIC].base,
                             memmap[NUCLEI_N_ECLIC].size, 51);
if( s->eclic != NULL)
{
    s->timer.eclic = s->eclic;
    s->timer.soft_irq = &(NUCLEI_ECLIC(s->eclic)->irqs[Internal_SysTimerSW_IRQn]);
    s->timer.timer_irq = &(NUCLEI_ECLIC(s->eclic)->irqs[Internal_SysTimer_IRQn]);

    s->uart0.irq = nuclei_eclic_get_irq(DEVICE(s->eclic), 22);
}
```

```
static void nuclei_mtimecmp_cb(void *opaque) {
    RISCVCPU *cpu = RISCV_CPU(qemu_get_cpu(0));
    CPURISCVState *env = &cpu->env;
    nuclei_eclic_systimer_cb(((RISCVCPU *)cpu)->env.eclic);
    timer_del(env->timer);
}
```

hw\intc\nuclei\_systimer.c

## SYSTIMER中

```
env->timer = timer_new_ns(QEMU_CLOCK_VIRTUAL,
                          &nuclei_mtimecmp_cb, cpu);
```

## timecmp对比: 将来触发

```
uint64_t next_ns = qemu_clock_get_ns(QEMU_CLOCK_VIRTUAL)+muldiv64(diff, NANoseconds_PER_SECOND,
s->timebase_freq);
timer_mod(env->timer, next_ns);
```

hw\intc\nuclei\_eclic.c

```
DeviceState *nuclei_eclic_create(hwaddr addr,
                                uint32_t aperture_size, uint32_t num_sources)
{
    DeviceState *dev = qdev_new(TYPE_NUCLEI_ECLIC);

    qdev_prop_set_uint32(dev, "aperture-size", aperture_size);
    qdev_prop_set_uint32(dev, "num-sources", num_sources);

    sysbus_realize_and_unref(SYS_BUS_DEVICE(dev), &error_fatal);
    sysbus_mmio_map(SYS_BUS_DEVICE(dev), 0, addr);
    return dev;
}
```

## 软件中断:

```
case NUCLEI_SYSTIMER_REG_MSIP:
    s->msip = value;
    if ((s->msip & 0x1) == 1) {
        qemu_set_irq(*(s->soft_irq), 1);
    } else {
        qemu_set_irq(*(s->soft_irq), 0);
    }
}
```



## NUCLEI ECLIC之测试

```
$qemu-system-riscv32 \  
-nographic -machine mcu_200t \  
-kernel ./hbird/demo_eclic.elf \  
-nodefaults -serial stdio
```

```
$qemu-system-riscv32 \  
-nographic -machine mcu_200t \  
-kernel ./hbird/demo_eclic2.elf \  
-nodefaults -serial stdio
```

```
Nuclei SDK Build Time: Apr 9 2021, 09:56:38  
Download Mode: ILM  
CPU Frequency 1182112 Hz  
Initialize timer and start timer interrupt periodically
```

```
[IN TIMER INTERRUPT]timer interrupt hit 0 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt will run when timer interrupt finished  
[IN TIMER INTERRUPT]timer interrupt end  
[IN SOFTWARE INTERRUPT]software interrupt hit 0 times  
[IN SOFTWARE INTERRUPT]software interrupt end
```

```
-----  
[IN TIMER INTERRUPT]timer interrupt hit 1 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt will run when timer interrupt finished  
[IN TIMER INTERRUPT]timer interrupt end  
[IN SOFTWARE INTERRUPT]software interrupt hit 1 times  
[IN SOFTWARE INTERRUPT]software interrupt end
```

```
-----  
[IN TIMER INTERRUPT]timer interrupt hit 2 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt  
[IN TIMER INTERRUPT]timer interrupt end  
[IN SOFTWARE INTERRUPT]software interrupt  
[IN SOFTWARE INTERRUPT]software interrupt end
```

```
-----  
[IN TIMER INTERRUPT]timer interrupt hit 14 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt will run during timer interrupt  
[IN SOFTWARE INTERRUPT]software interrupt hit 14 times  
[IN SOFTWARE INTERRUPT]software interrupt end  
[IN TIMER INTERRUPT]timer interrupt end
```

```
-----  
[IN TIMER INTERRUPT]timer interrupt hit 15 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt will run during timer interrupt  
[IN SOFTWARE INTERRUPT]software interrupt hit 15 times  
[IN SOFTWARE INTERRUPT]software interrupt end  
[IN TIMER INTERRUPT]timer interrupt end
```

```
-----  
[IN TIMER INTERRUPT]timer interrupt hit 16 times  
[IN TIMER INTERRUPT]trigger software interrupt  
[IN TIMER INTERRUPT]software interrupt will run during timer interrupt  
[IN SOFTWARE INTERRUPT]software interrupt hit 16 times  
[IN SOFTWARE INTERRUPT]software interrupt end  
[IN TIMER INTERRUPT]timer interrupt end
```

## 下节课内容:

### 中断虚拟化

- ECLIC 与 CLIC
- SYSTIMER 与 CLINT
- PLIC介绍
- ECLIC在UART中的使用

# 谢谢

wangjunqiang@iscas.ac.cn